


Lecture 8: Back to inheritance!

- Last lecture we extended our Matrix class to use templating
- That allows us to build our Matrix class using doubles, ints or other numeric types that have a `*` and `[]` operators defined
- We're going to further extend our Matrix "library" through inheritance

Sparse matrices

- Our existing Matrix class stored dense matrices
- We spoke a little about different types of “sparse” matrices
- These are matrices where some entries are zero




Banded matrices

- For when non-zero entries are a set amount away from the diagonal
- Very common way to store matrices that come from (local) spatial discretisations of PDEs (e.g., see ACSE-3 Lecture 7)

$$\begin{pmatrix} X & X & X & \cdot & \cdot & \cdot & \cdot \\ X & X & \cdot & X & X & \cdot & \cdot \\ X & \cdot & X & \cdot & X & \cdot & \cdot \\ \cdot & X & \cdot & X & \cdot & X & \cdot \\ \cdot & X & X & \cdot & X & X & X \\ \cdot & \cdot & \cdot & X & X & X & \cdot \\ \cdot & \cdot & \cdot & \cdot & X & \cdot & X \end{pmatrix}$$



Banded matrices

- How would we store a banded matrix?
- How would we do a matrix-matrix product in that format?



Algorithms for changing sparsity

- The sparsity of a matrix is a consequence of the ordering of the unknowns
- The solution is still the same
- Algorithms exist to reorder the unknowns in desirable ways
- E.g., RCM (reverse Cuthill-McGee) to reduce bandwidth



Symmetric matrices

- Symmetry can be exploited to reduce the storage, either in dense, banded or CSR storage of a symmetric matrix
- Not quite half the storage (given the diagonal)

$$A = \begin{pmatrix} 1 & -1 & * & -3 & * \\ -1 & 5 & * & * & * \\ * & * & 4 & 6 & 4 \\ -3 & * & 6 & 7 & * \\ * & * & 4 & * & -5 \end{pmatrix}$$

CSR matrices

- Known as compressed sparse row
- A store the values, IA the index into each row, JA the column for each non-zero

For example, the matrix

$$\begin{pmatrix} 0 & 0 & 0 & 0 \\ 5 & 8 & 0 & 0 \\ 0 & 0 & 3 & 0 \\ 0 & 6 & 0 & 0 \end{pmatrix}$$

is a 4×4 matrix with 4 nonzero elements, hence

$$\begin{aligned} A &= [5 \ 8 \ 3 \ 6] \\ IA &= [0 \ 0 \ 2 \ 3 \ 4] \\ JA &= [0 \ 1 \ 2 \ 1] \end{aligned}$$

Time to write some code!

- Get your Matrix.h and Matrix.cpp from last lecture
- Now it's time to examine some design decisions we made when we created Matrix.h and Matrix.cpp
- Why was “preallocated” private?
- Why did we make matmatMult virtual?
- We are going to create CSRMatrix.h and CSRMatrix.cpp

Matrix-vector product

- How easily vectorised do you think the matrix-vector product we wrote is?

```
output[i] += this->values[val_index] *  
    input[this->col_index[val_index]];
```

CSRMatrix

- Let's think about our `matmatMult(mat_right, output)`
- How do we know the sparsity of our output CSRMatrix
- Often sparse libraries will compute a “symbolic `matmatMult`”
- This just loops over the sparsity as if it were doing a matrix matrix product
- $A = B * C$
- If there aren't entries in the corresponding sparsity of B and C, there won't be a non-zero entry in A
- If the sparsity doesn't change between `matmatMults`, you don't have to redo the symbolic

Existing libraries

- There are lots of sparse matrix libraries (that also include lots that already exist)
- Rule number 1: Never write code if someone clever has already done the hard work for you!
- PETSc (Portable Extensible Toolkit for Scientific Computation) (<https://www.mcs.anl.gov/petsc/>)
- Trilinos (github.com/trilinos/Trilinos)
- Intel MKL

PETSc

- Heavily used and supported
- Lots of different dense/sparse Matrix classes
- Lots of different linear solvers (including direct/iterative solvers)
- For example look at all the different Matrix methods they provide:
- The serial PETSc CSR class is called MATAIJ

<https://www.mcs.anl.gov/petsc/petsc-current/docs/manualpages/Mat/index.html>

PETSc methods to consider:

- MatCreateAIJ
- MatSetValue
- MatBeginAssembly
- MatEndAssembly
- MatMatMultSymbolic
- MatMatMultNumeric
- We can take inspiration from these methods
- “Under the hood” a lot of these methods make calls to BLAS/LAPACK!

Homework

- Rewrite our `CSRMatrix::printMatrix` to print out in a 2D pattern
- Build a `setvalues()` function for our `CSRMatrix` inspired by PETSc that takes in i, j , and a value, and sets the value in that position (maybe include something like “Assemble”)
- ADVANCED - Build a symbolic matrix-matrix product for our `CSRMatrix`
- SUPER ADVANCED – Build our `matMatMult` for `CSRMatrix`, which calls the symbolic matrix-matrix product to know how many nnzs to allocate for the output
- ULTRA ADVANCED – Try using PETSc, either install/link and call from C++, or use PETSc4Py on a Jupyter notebook